

Docker

Before Docker :

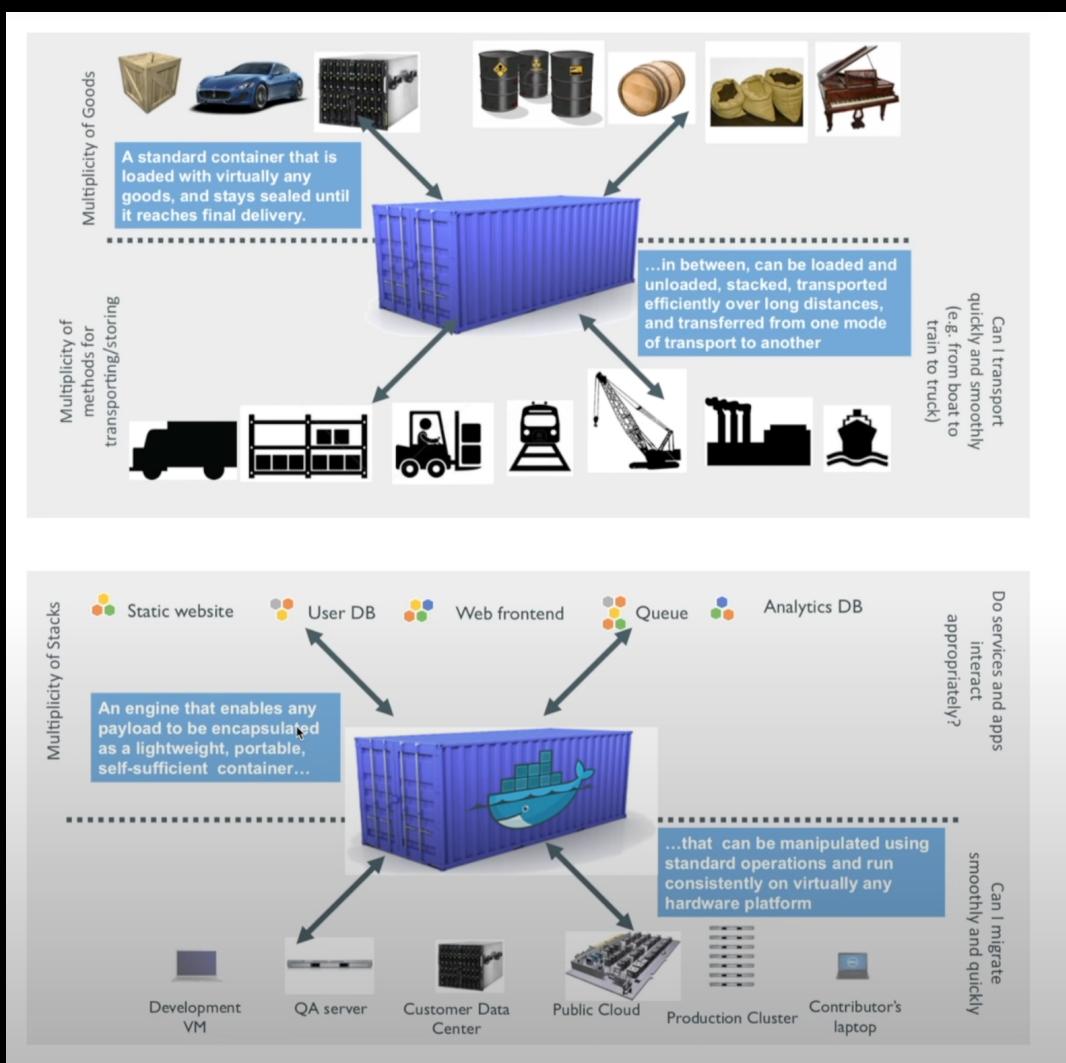
- Initially only 1 app used to run on one server

The diagram shows a rectangular box labeled 'Server' containing a smaller box labeled 'APP' with an arrow pointing to it.
- So, if you want to add an app, you basically have to add a server
- VMWARE Solves this problem by introducing Virtual Machines.
- VM's allows us to run multiple apps on same server
- Limitations: They require their own OS.
- Migration is time consuming

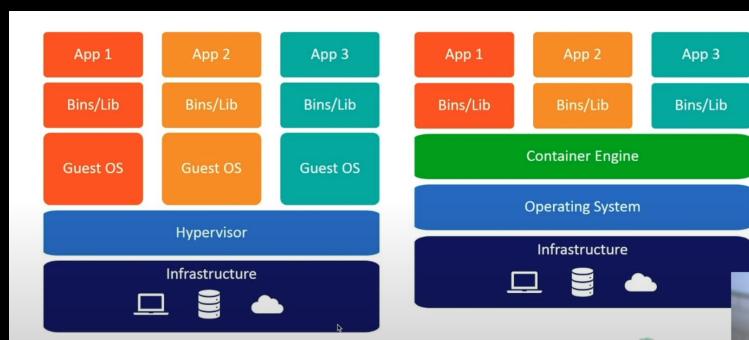
→ "It's work on my machine"

When you try to send your code
to run on your friend's Comp but
unable to due to module version &
so on

Containers:



- Like In The above example we you want transport to diff Country how all the Box, Piano, Car are put In 1 Box & Ship?
- Similarly you put the Code, dependency, everything required to run the code, you put In Box 1 give to friend now if will make live of will run In your friend computer.



VM

Containers

- * Hypervisor Is used to Create multiple VM's on the host OS & It manages the VM.
- Every OS will have dedicated hardware

→ Containers have one OS only & you have a Container Engine.

→ Running Docker

→ Docker windows Container won't work on docker Linux. So, you have to Refill image.

Windows: Docker desktop + WSL

MAC: VMs & Docker Desktop

Linux: As It Is

What Is Docker?

Docker is a Container platform, that allows us to build, test & deploy apps quickly.

- So, far we know Docker is helping us run in an isolated environment.
- It means container doesn't know what's happening outside even in the OS.

1. Runtime :

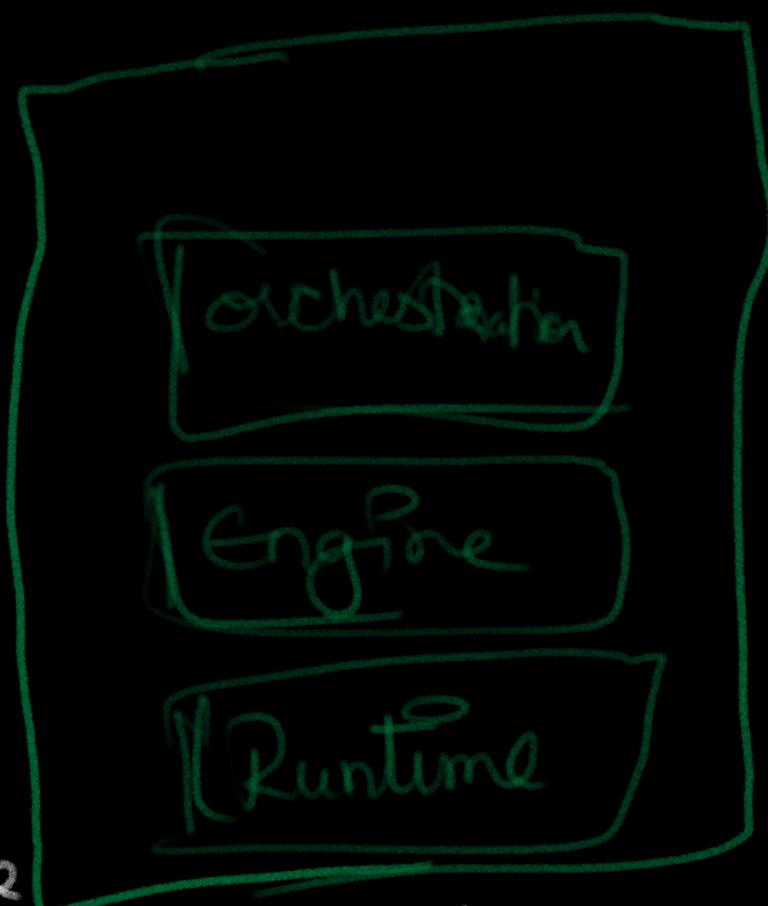
It allows us to start & stop the containers

- RunC

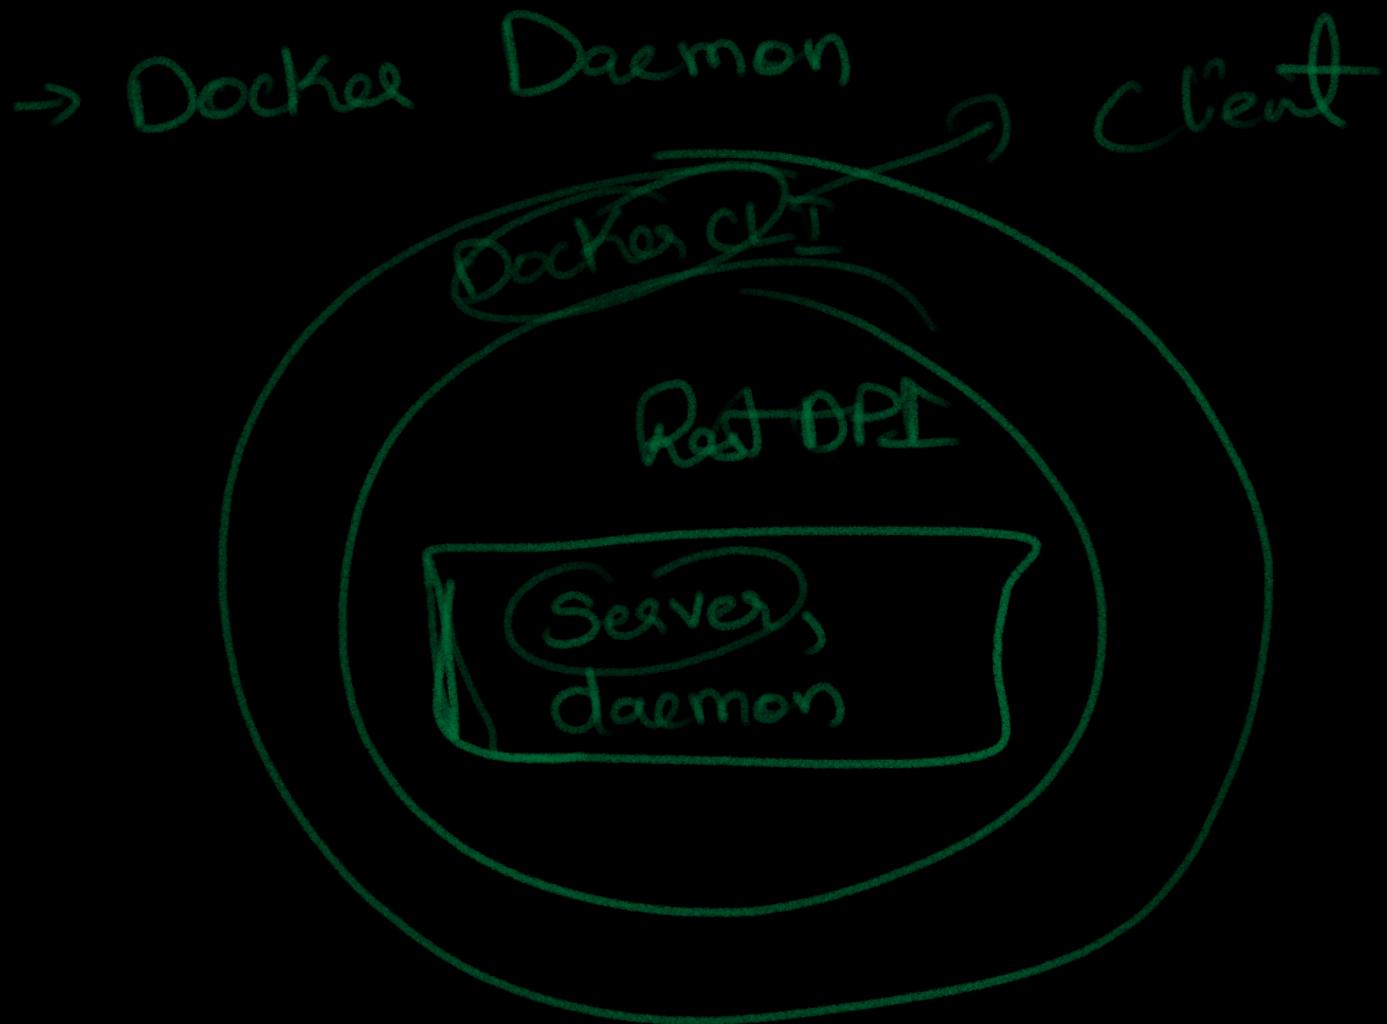
- ↳ low level runtime
- ↳ role is to work with OS, start & stop the containers.

- Containerd

- ↳ managing sync 2 Containers
- ↳ pulling the Images, get data in network.



2. Docker Engine



→ So, using docker CLI & Rest API,
can make an API call to daemon,
daemon will work with Container
runTime & def.

→ Uses Client-Server Architecture

3. Orchestration

Let's say an app is running 10K containers & you want to scale it to 20K.

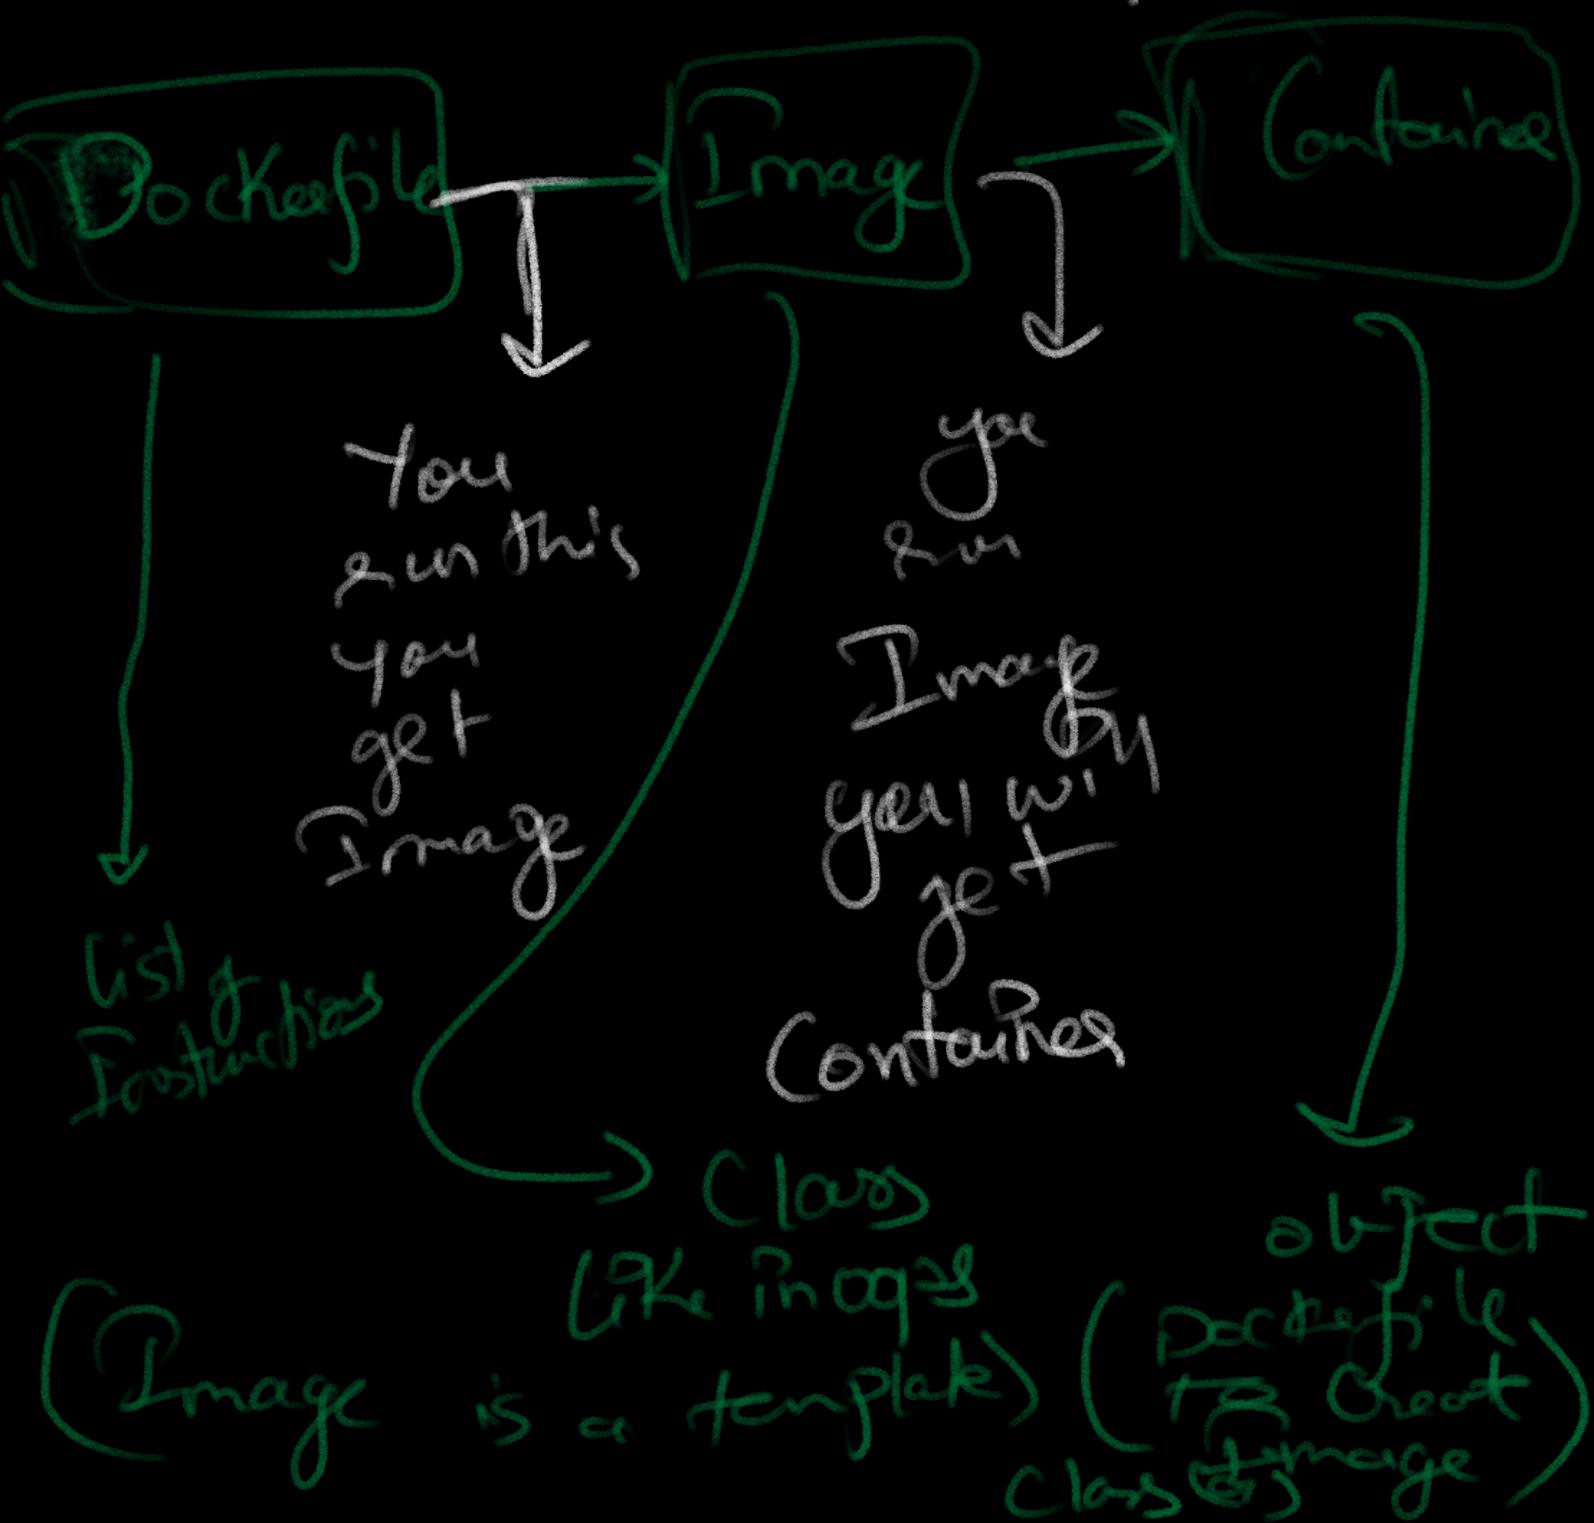
- Assume 500 containers got killed & you are going to restart
- It allows us to manage containers

Container Image

→ Let's say you are making a dish eating something ex: noodles & your friend lives in diff city & you can't just ship noodle soup in box to your friend when you are far away. It doesn't make sense to ship the running App

- Instead we can share the Instructions / Recipe.
- Similarly, You can send the entire Instruction with the code for website & tell your friend run it.
- This Instruction part, you follow recipe the dish will be created dish will be the actual app running in servers in containers.
- The thing that leads us to create the dish is known as Docker / Container Image
- So, if you want to share your app make an docker image & send it & the others can run.

- Contains a running instance of an Image.
- You can create docker Images using Dockerfile
 - ↳ Recipe.



→ There are many runtimes, so, how can we choose then they come up with an initiative known as Open Container Initiative

→ OCI contains 2 specs:

1. Runtime Spec
2. Image Spec

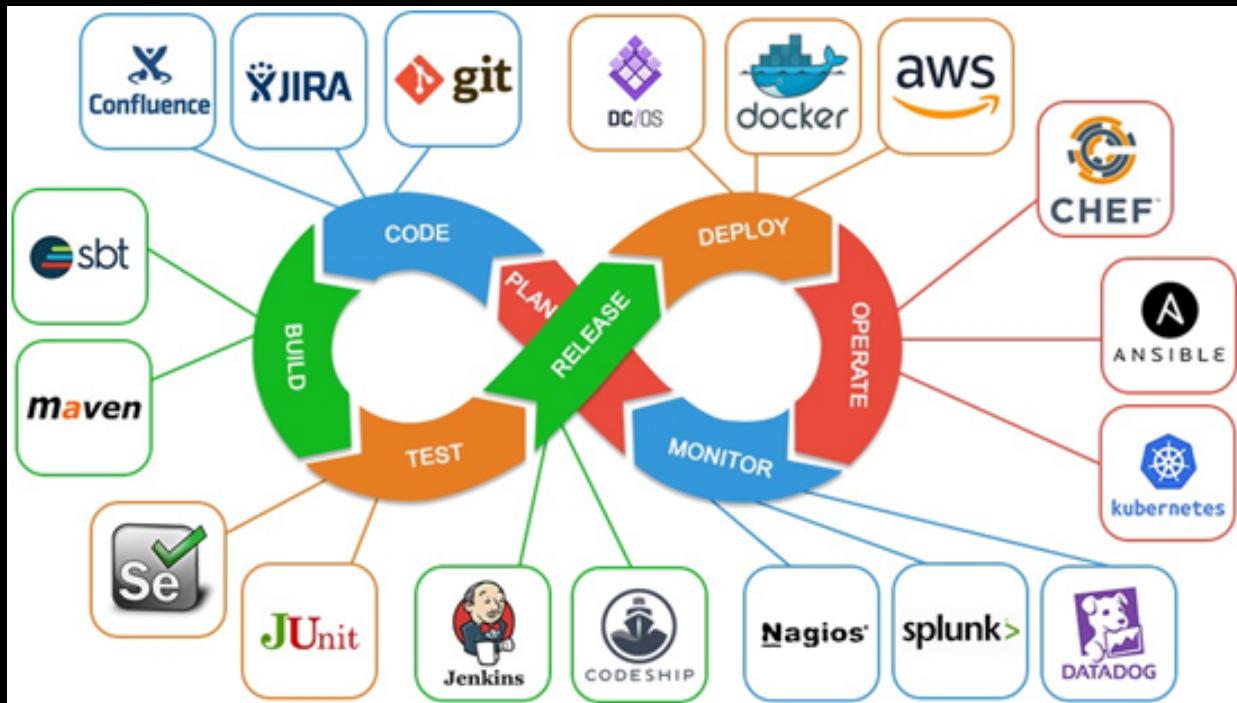


What are the various Container Runtimes?

Some of the popular container runtimes are below ones:

- **containerd**: A CNCF project, it manages the complete container lifecycle of its host system that includes image management, storage and container lifecycle, supervision, execution and networking.
- **lxc**: LXC provides OS level virtualization through a virtual environment that has its own process and network space, it uses linux cgroups and namespaces to provide the isolation.
- **runc**: runc is a CLI tool for spawning and running containers according to the OCI specification. It was developed by Docker Inc and donated to OCI as the first OCI runtime-spec compliant reference implementation.
- **cri-o**: CRI-O is an implementation of the Kubernetes CRI (Container Runtime Interface) to enable using OCI (Open Container Initiative) compatible runtimes. It is a lightweight alternative to using Docker as the runtime for Kubernetes.
- **rkt**: rkt is an application container engine developed by

Dev OPS



⇒ Development & operations

Dev OB

APP
Dockerfile

Image
Container

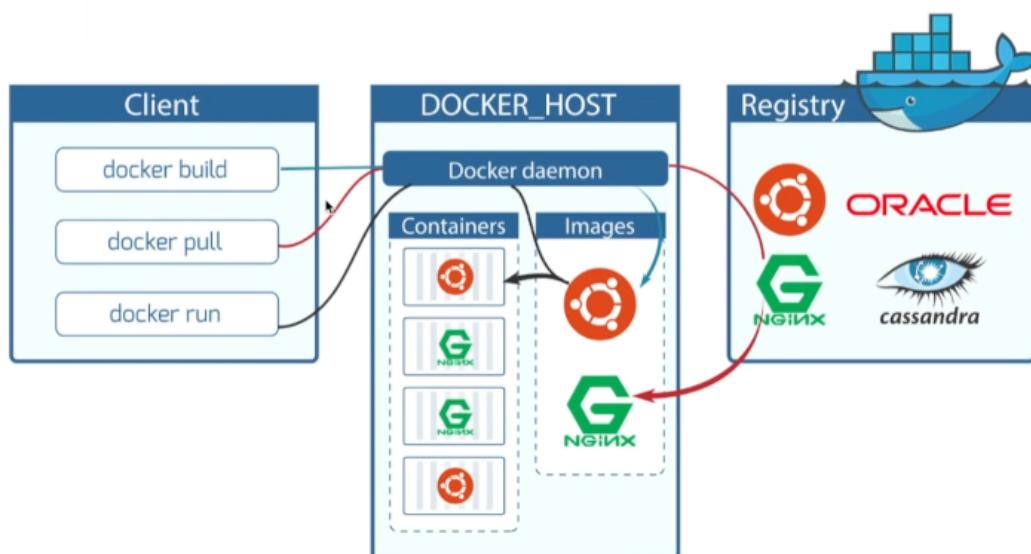
Download
Image,
run Image,
Creating
Image.

- If you are not using containers the operation team need to handle hosting & deployment.
 - Simple, if you are using containers you can use node, python, swift anything you want as long as you are creating Images.
-

Docker CLI

\$ Docker run hello-world → Image runs
↳ run an Image to create a Container.

DOCKER COMPONENTS



→ Here when you enter docker run
it will talk to docker daemon,
it will check whether the
daemon will check whether the
image is present in local system
If not it will go the registry
Docker hub (Image library)
& it will download from there &
runs image & forms the container.

→ As you know we need OS to run any APP like mongoDB but since Container doesn't interact with outside world then where is the mongo dB running?

When we talk about docker Image they also contain operating system files. It contains smaller version of OS in the BoxImage

→ If you don't want to run you can directly download using \$ docker pull

→ To view the running containers
\$ docker ps

→ \$ docker Images to view all the Images

→ docker Container ls

Shows all the running containers

```
docker run -it ubuntu:16.04
root@90647413e3b2:/# ls
bin  boot  dev  etc  home  lib  media  mnt  opt  pr
oc  root  run  sbin  srv  sys  tmp  usr  var
root@90647413e3b2:/# ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT
START      TIME COMMAND
root        1  0.0  0.0    3544  2560 pts/0      Ss
04:47     0:00 /bin/bash
root       12  0.0  0.0    5472  2304 pts/0      R+
04:48     0:00 ps aux
root@90647413e3b2:/# ls
bin  dev  home  media  opt  root  sbin  sys  usr
boot  etc  lib  mnt  proc  run  srv  tmp  var
root@90647413e3b2:/# touch hey.txt
root@90647413e3b2:/# ls
bin  etc  lib  opt  run  sys  var
boot  hey.txt  media  proc  sbin  tmp
dev  home  mnt  _root  srv  usr
```

→ On a shell, running Ubuntu
& Created a new file `hey.txt`
&

→ Now using the docker
Container ls to know Container
& we did below

```
base ~ (0.057s)
docker container ls

CONTAINER ID        IMAGE       COMMAND      CREATED      NAMES
90647413e3b2        ubuntu:16.04 "/bin/bash"   About a minute ago    About a vigorous_ellis
```

```
base ~
docker container exec -it 90647413e3b2 bash

root@90647413e3b2:/# ls
bin  etc  lib  opt  run  sys  var
boot  hey.txt  media  proc  sbin  tmp
dev  home  mnt  root  srv  usr
root@90647413e3b2:/#
```

- It means the Bash shell should be attached to this running container
- This command won't work if the container is not running
- As you can see `hey.txt` is present here as well.
- Both shells are attached to same container

→ You can also copy Container ID
if we use `$ docker stop` IP
to stop running

→ `$ docker ps -a` → Shows all the stopped
Containers

→ To remove Container `$ docker rm ID`

→ `$ docker inspect` (e.g.: Ubuntu ID)
gives all info about Container

↳ Image state, path, ID,
when created, running or not,
process ID, network info.

→ `$ docker logs ID` to check
history

→ \$ Docker Container prune -f
means deletes all the containers that
are stopped.

Prune - deletes all containers that stopped
-f - don't ask again, just do it

base ~ (0.458s)

docker run -d alpine ping www.apple.com

616cf8fc5bdc16a55d638074bc736c13a538bbe7a9ffc0ed95d
a75e0923d95c9

- ↳ -d means detached mode
- ↳ runs in background,
- ↳ returned Id of container

→ Alpine is image based on
alpine linux, since all Images
are based on Linux Kernels
you can run Linux Commands like
ping.

base ~ (0.489s)

docker run ubuntu echo hey

hey

→ ↴ whatever you put outside
like this will be running it
on linux terminal

base ~ (0.062s)

docker logs 616c

```
PING www.apple.com (23.48.49.37): 56 data bytes
64 bytes from 23.48.49.37: seq=0 ttl=63 time=122.146 ms
64 bytes from 23.48.49.37: seq=1 ttl=63 time=32.962 ms
64 bytes from 23.48.49.37: seq=2 ttl=63 time=32.531 ms
64 bytes from 23.48.49.37: seq=3 ttl=63 time=30.907 ms
64 bytes from 23.48.49.37: seq=4 ttl=63 time=32.498 ms
64 bytes from 23.48.49.37: seq=5 ttl=63 time=37.084 ms
64 bytes from 23.48.49.37: seq=6 ttl=63 time=38.257 ms
64 bytes from 23.48.49.37: seq=7 ttl=63 time=38.093 ms
64 bytes from 23.48.49.37: seq=8 ttl=63 time=34.407 ms
64 bytes from 23.48.49.37: seq=9 ttl=63 time=40.316 ms
64 bytes from 23.48.49.37: seq=10 ttl=63 time=84.799 ms
64 bytes from 23.48.49.37: seq=11 ttl=63 time=105.360 ms
64 bytes from 23.48.49.37: seq=12 ttl=63 time=31.836 ms
64 bytes from 23.48.49.37: seq=13 ttl=63 time=31.997 ms
64 bytes from 23.48.49.37: seq=14 ttl=63 time=29.808 ms
64 bytes from 23.48.49.37: seq=15 ttl=63 time=31.567 ms
64 bytes from 23.48.49.37: seq=16 ttl=63 time=33.922 ms
64 bytes from 23.48.49.37: seq=17 ttl=63 time=39.026 ms
```

→ first 4 digits will suffice

```
base ~ (0.069s)
docker logs --since 5s 616c

64 bytes from 23.48.49.37: seq=508 ttl=63 time=38.362 ms
64 bytes from 23.48.49.37: seq=509 ttl=63 time=37.803 ms
64 bytes from 23.48.49.37: seq=510 ttl=63 time=36.373 ms
64 bytes from 23.48.49.37: seq=511 ttl=63 time=33.058 ms
64 bytes from 23.48.49.37: seq=512 ttl=63 time=37.693 ms
```

↳ Here it gives logs for only last 5 seconds

```
base ~ (0.069s)
docker rmi alpine -f

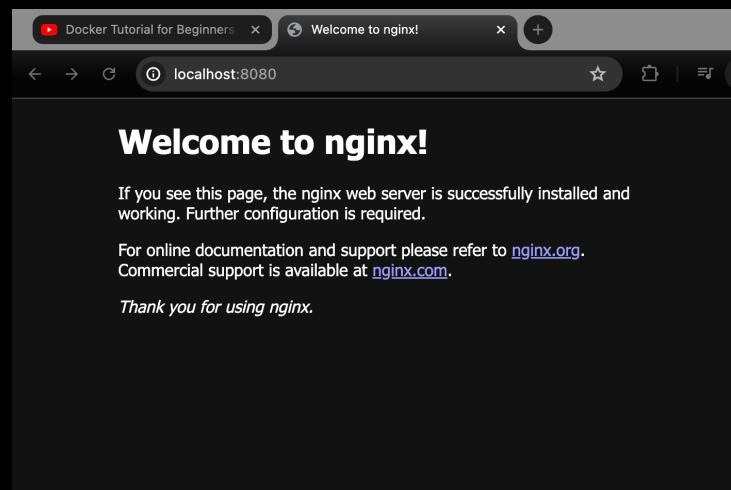
Untagged: alpine:latest
Untagged: alpine@sha256:b89d9c93e9ed3597455c90a0b88a8bbb5cb7188438f70953fede212a0c4394e0
Deleted: sha256:092561eea88f9f28654bbade209576f5f93efeb8c7ba66a07ac2033c0ddc8ae7
```

↳ Force Remove Image

```
base ~ (0.363s)
docker run -d -p 8080:80 nginx
24b5fda024cd7e28b1edcc103dd9ebac9108d7e90e7732c0a015c6944ebdd1b1
```

↳ p is taking 2 arguments the ports you want to expose it to & port of Container

8080 → exposed port
localhost:8080 → port of host machine



60 → port of Container. (default port of nginx → 80)

→ So, for accessing Container locally we do as above In the Image, It does port forwarding

```
base ~
```

```
docker run -it ubuntu
```

```
root@b4bc0a4a72:/# touch names.txt
```

```
root@b4bc0a4a72:/# echo " Hey my name is Akash" > names.txt
```

```
root@b4bc0a4a72:/# cat names.txt
```

```
Hey my name is Akash
```

```
root@b4bc0a4a72:/# █
```

↳ If In a new shell again you run the above command names.txt will not be there because It will be creating a new Container

```
base ~ (3m 4.16s)
docker run -it ubuntu
root@b4bc0a4a72:/# touch names.txt
root@b4bc0a4a72:/# echo " Hey my name is Akash" >
> names.txt
root@b4bc0a4a72:/# cat names.txt
Hey my name is Akash
root@b4bc0a4a72:/# exit
exit
```

```
base ~ (0.061s)
docker ps -a
CONTAINER ID   IMAGE      COMMAND
CREATED        STATUS      PORTS
S   NAMES
b4bc0a4a72    ubuntu     "/bin/bash"
3 minutes ago  Exited (0) 18 seconds ago
determined_no
cee0b75ba504  ubuntu     "echo hey"
21 minutes ago Exited (0) 21 minutes ago
brave_darwin
0f2f12850381  092561ee88f  "ping www.apple.com"
27 minutes ago Exited (0) 26 minutes ago
pedantic_morse
```

```
base ~ (0.239s)
docker start b4bc0a4a72
b4bc0a4a72
```

```
base ~
docker exec -it b4bc0a4a72 bash
root@b4bc0a4a72:/# ls
bin  etc  media  opt  run  sys  var
boot  home  mnt  proc  sbin  tmp
dev  lib  names.txt  root  srv  usr
root@b4bc0a4a72:/#
```

```
base ~ (0.056s)
docker ps
CONTAINER ID   IMAGE      COMMAND
CREATED        STATUS      PORTS
S   NAMES
b4bc0a4a72    ubuntu     "/bin/bash"  5 seconds ag
o Up 4 seconds
determined_no
```

```
base ~ (0.066s)
docker ps
CONTAINER ID   IMAGE      COMMAND
CREATED        STATUS      PORTS
S   NAMES
b4bc0a4a72    ubuntu     "/bin/bash"  3 minutes ag
o Up 14 seconds
determined_no
```

```
base ~
Try typing natural language instead of a command
```

↳ once you edit container stops
So using docker ps (-a) docker ps -a
find the container ID & then
use # Docker Start ContainerID &
then use docker exec -it id bash
to fake terminal

```
base ~ (0.13s)
docker commit -m "added names.txt file" b4bc0a4a72 names_ubuntu:1.01
sha256:47109159790163093984ac6835ed1bc625c22684a106b17d4ae4cc98dac31c7f
```

```
base ~ (0.063s)
docker images
```

REPOSITORY	IMAGE ID	CREATED	TAG	SIZE
names_ubuntu	471091597901	5 seconds ago	1.01	101MB

↳ By using this Command

Container ID
↳ Image name, Version

```
base ~ (31.286s)
docker run -it names_ubuntu:1.01
root@e67286863357:/# ls
bin  boot  dev  etc  home  lib  media  mnt  names.txt  opt  proc  root  run  sbin  srv  sys  tmp  usr  var
root@e67286863357:/# exit
exit
```

```
base ~ (0.063s)
docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
TS	names_ubuntu:1.01	/bin/bash"	43 seconds ago	Exited (0) 12 seconds ago		
b4bc0a4a72	ubuntu	/bin/bash"	13 minutes ago	Exited (137) 3 minutes ago		
cee0b75ba504	ubuntu	"echo hey"	31 minutes ago	Exited (0) 31 minutes ago		
0f2f12850381	092561eea88f	"ping www.apple.com"	37 minutes ago	Exited (0) 36 minutes ago		
	pedantic_morse					

```
base ~ (0.052s)
docker stop e67286863357
e67286863357
```

```
base ~ (0.061s)
docker images -q
471091597901
bd1f209b12c1
26af66f0ce61
60e5a9498560
23e08b4032bf
d5e32897ae01
3970bf581bd2
d49959718993
a3b99b6de3e7
5784497e97da
f1f044275595
89424ddaf8a6
6f8fb6bbe9f
7c5ce99f18b3
b897f7f0612e
a5d57a9e6115
bec7a14748a1
2d9e1b1f31ce
695e39d4c41b
```

↓
As you can see when we run names.txt exist

- ↳ It only gives ID of all Images
- - ↳ removes all the Images
 - ↳ Also you can you -f to force it
- But remember If won't delete running containers

```
→ ~ docker run ubuntu:16.04
Unable to find image 'ubuntu:16.04' locally
16.04: Pulling from library/ubuntu
828b35a09f0b: Downloading 2.971MB/41.24MB
238e9b8fdf46: Download complete
1d5b1b491959: Download complete
269a6c6175ba: Download complete
```

- These 4 things are called layers
- Images are built in layers.
- Each layer is immutable file, A collection of files.
- Ex:- Here Ubuntu Image is collection of 4 files

why?

- Since, Everything is based on Linux Kernel there may be a possibility if you downloading another Image let's say mongoDB It may give you some files that are also being used by Ubuntu for ex
- Since you downloaded Ubuntu Image, It does not make sense to download same files for other Image.
- So, Images are built in layers If Every Images has a hash Value calculated by SHA 256

base ~ (0.069s)

docker images -q --no-trunc

```
sha256:47109159790163093984ac6835ed1bc625c22684a106b17d4ae4cc98dac31c7f
sha256:bd1f209b12c1d2acfac2c4d063c093a39499e8db44475073cb7ab427b09012d5
sha256:26af66f0ce61786c7c80a4c3df3fc82449ba62eb649d2a456841697910e1e988
sha256:60e5a9498560419431c1dd00c4001c549656390999976d2bbef46afc0a655bc
sha256:23e08b4032bf4d5865f8d37fdb121697e6eb335fa26e08683d1bd26cb3a72c1
sha256:d5e32897ae0120e97b2e5fb7d8cabc3c71b80db2e8ba50f9047cc16b4abe638b
sha256:3970bf581bd22fd83c054b77e726be5a8b716006990eaaf88dfcf6e73e1ab68
sha256:a3b99b6de3e7668ded16f3fbff3d07ffd9ac61853671d684a7f682932e39b8f
sha256:d49959718993a2a576792eefa24b3ae217882bf190899db3bfa9cf3478826676
sha256:5784497e97da1d5a239a59b545575d62fd2249b8291931d5fb8f22e37f7ba8b
sha256:f1f044275595b8324130d12b8d97becf965360f899051ebff2486eeb4b9d4541
sha256:89424ddf8a6a51b50739059458643e452092bc8d92fbf888cbee2656c4af
sha256:6f8fb6bbeb9f534c1acb64cdccfedbe0d8994188b4a8c73b2877ca910e929896
sha256:7c5ce99f18b31c735359a52d165caf562f1f5b99983d43c71385bff7877bb747
sha256:b897f7f0612ef880127b7624e6f649b30af036d9a2709a1f8949298cfa622b3e
sha256:a5d57a9e6115b2505c8c5a84fe05b40f751f561e18153bdb1e6f4ddaf03d7d0e
sha256:bec7a14748a14a45b4b43e484cb8f07cc329eae1a1742af5dc93a05780b8aec4
sha256:2d9e1b1f31ce77ba4425293c3585b4885593b86cccb9e6856e89a8749c841d6e
sha256:695e39d4c41b769cd71053d787aa47785ed3d7231d5c3056faa03b4c8c085b2f
sha256:9572b90f89886750618efaade1613e782e421156a23fbe4aa38c9a5fb15d6601
sha256:ffb64c97cde8b9f1891f7a72609d4e691c4671dc4aa83084fc7b5774958d827de
sha256:11ceee7cdc57225711b8382e1965974bbb259de14a9f5f7d6b9f161ced50a10a
sha256:c8405f018c81730ffbf948a84e63d825d259df975b5b9500fd6426ea48fb27
sha256:2bceedbd3273e5cc709fc640b29d04fb7dde67b9261dcdf3f69d4a02efca7219
sha256:32eb28f276a5415919c40d989d9a5d49852dcbb8d18d8b0f4510102f9a0c1d12
sha256:96e2c08ad32ed230afea843bed0bebdfc0bc379775f12615466645dac97a2ecf
sha256:6241433e83b5a5599b4d5c54927bcc77d330f52504cedb7eb7685b9b5d6bce4
sha256:ce59d6220a4de14dfa902c338d6025f55bfe27495b069fd37993082bf521370f
sha256:21efaf670a9c106c4db8748e3ff5e2b7cc95d824ee712522954f0c1d965412f2
sha256:499fa48260b8bd955d5891d6a1dc82d6e2ce7ca908b72ed66ee4a9939018ed9b
sha256:79f8d13ae8b8839cadfb2f83416935f5184206d386028e2d1263577f0ab3620b
sha256:789de7d5f83dc4ef225bcd637e49dd5edf4e22190546246045a513df6218a1
sha256:2437cf762177702dec2dfe99a09c37427a15af6d9a57c456b65352667c223d93
sha256:3750dfec169f630ab8042ca29bca08cbf436ccf1a74ae6f2778f31d0235651d5
sha256:ee301c921b8aadc002973b2e0c3da17d701dc994b606769a7e6aaa100b81d44
sha256:d125c6a1fe22504f84552b7eb11e353e88691875c18caf68847843892b190ccc
sha256:829e9de338bd5fd3f16f68f83a9fb288fb8453e881e5d5cf0f6f2ff72b43e
sha256:c027a58fa0bb1f8ff35555e068fc8d1a67173782978e25d0450bac0eb5848cd7
sha256:8a85ac985999aae3427246efaa562cc8d6ec4911721b245b71958a8aef693e1d
```

base ~ (0.063s)

docker images

REPOSITORY	IMAGE ID	CREATED	TAG
names_ubuntu	471091597901	21 minutes ago	1.01
mongo	bd1f209b12c1	3 days ago	latest

You can see
Image ID is
equal to the
1st few hashed
Value.

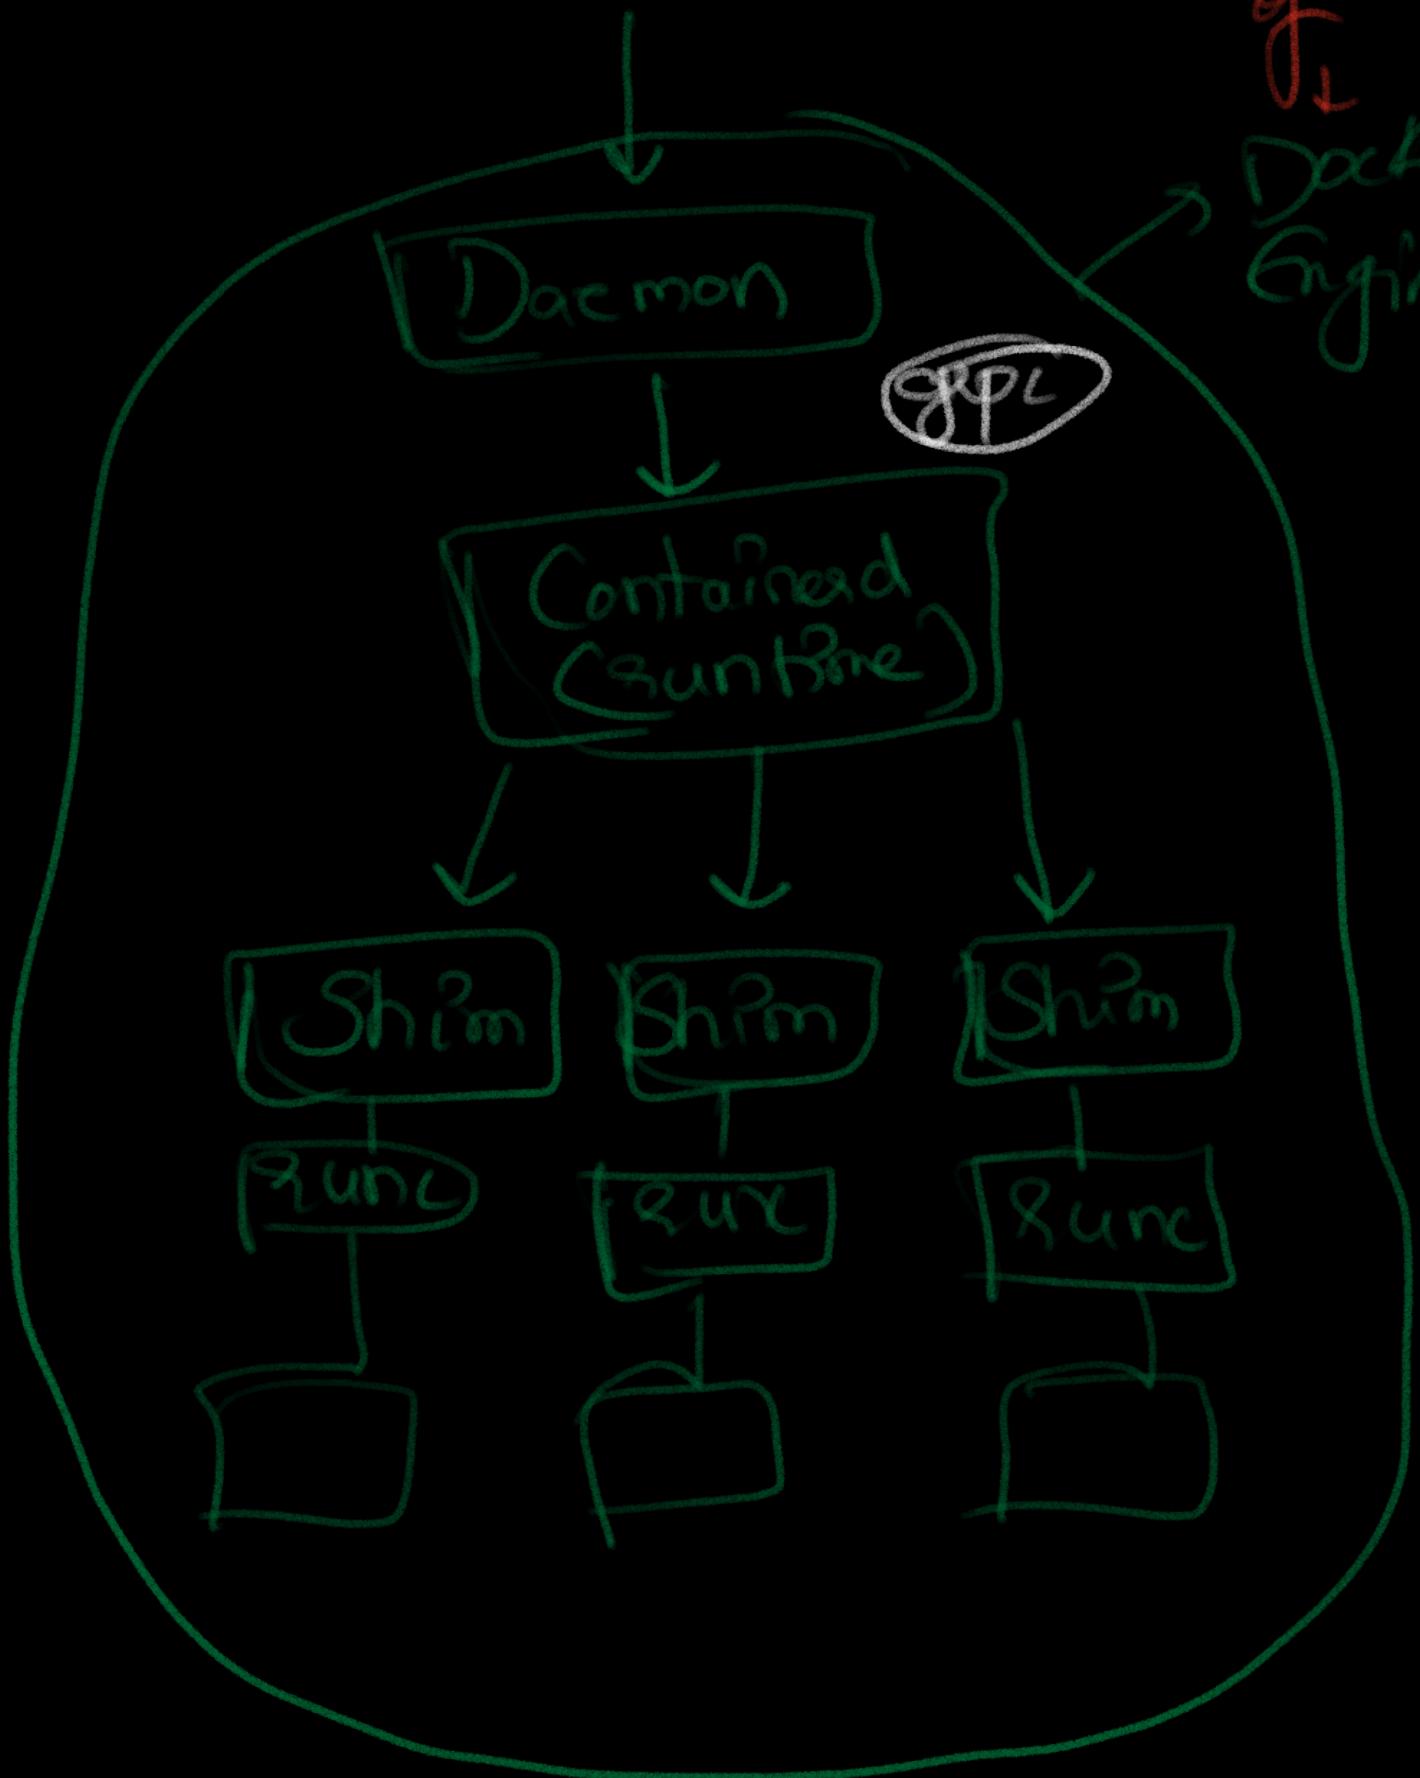
→ Also using **# docker Inspect**, you
can check the layers.

→ If there is a particular layer
of image that is already downloaded
on your system, it will not download
again. That's how it makes it fast

Docker client

Architecture

of
Docker
Engine



- When run the `# Docker run` command docker client, it talks to docker daemon talk to container using CRUOperations API's over gRPC (it's a protocol used by clients & servers to communicate basically)
- So, Containerd is not the one creating containers. RunC does that.
- Containerd will take the Image & pass it RunC & ask it to create a Container then RunC will talk to your OS kernels & pull resources & create a Container.
- Once the Daemon is stopped, all the containers running will also be stopped.

- > If you want to update your daemon than it will Impact running containers.
 - > So, Once the Container is Created RunC will be removed, RunC role is just deleting, starting Container.
 - > Shim takes over RunC once it goes, these are known as daemon less containers.
 - > Client $\xrightarrow{\text{HTTP(2375)}}$ Server
- TLS - Transport Layer Security
you can allow only which they exists.